

Which Questions we should have

Jun Kim , Senior Solution Architect &
Database Specialist



Agenda

1. Polyglot persistence & Multi-model Database
2. NoSQL Vs RDBMS
3. Considerations & Questions ?
4. MongoDB is ..
5. Proof of Points

Polyglot persistence Vs Multi-model database

Polyglot persistence is the concept of using different data storage technologies to handle different data storage needs within a given software application.

Most database management systems are organized around a single data model that determines how data can be organized, stored, and manipulated. In contrast, a **multi-model database** is designed to support multiple data models against a single, integrated backend. Document, graph, relational, and key-value models are examples of data models that may be supported by a multi-model database.

RDBMSs Vs NoSQLs

How To Use in the Market



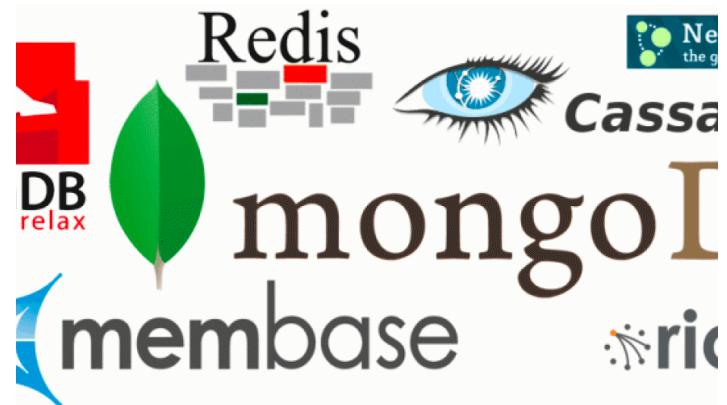
ORACLE®

PostgreSQL

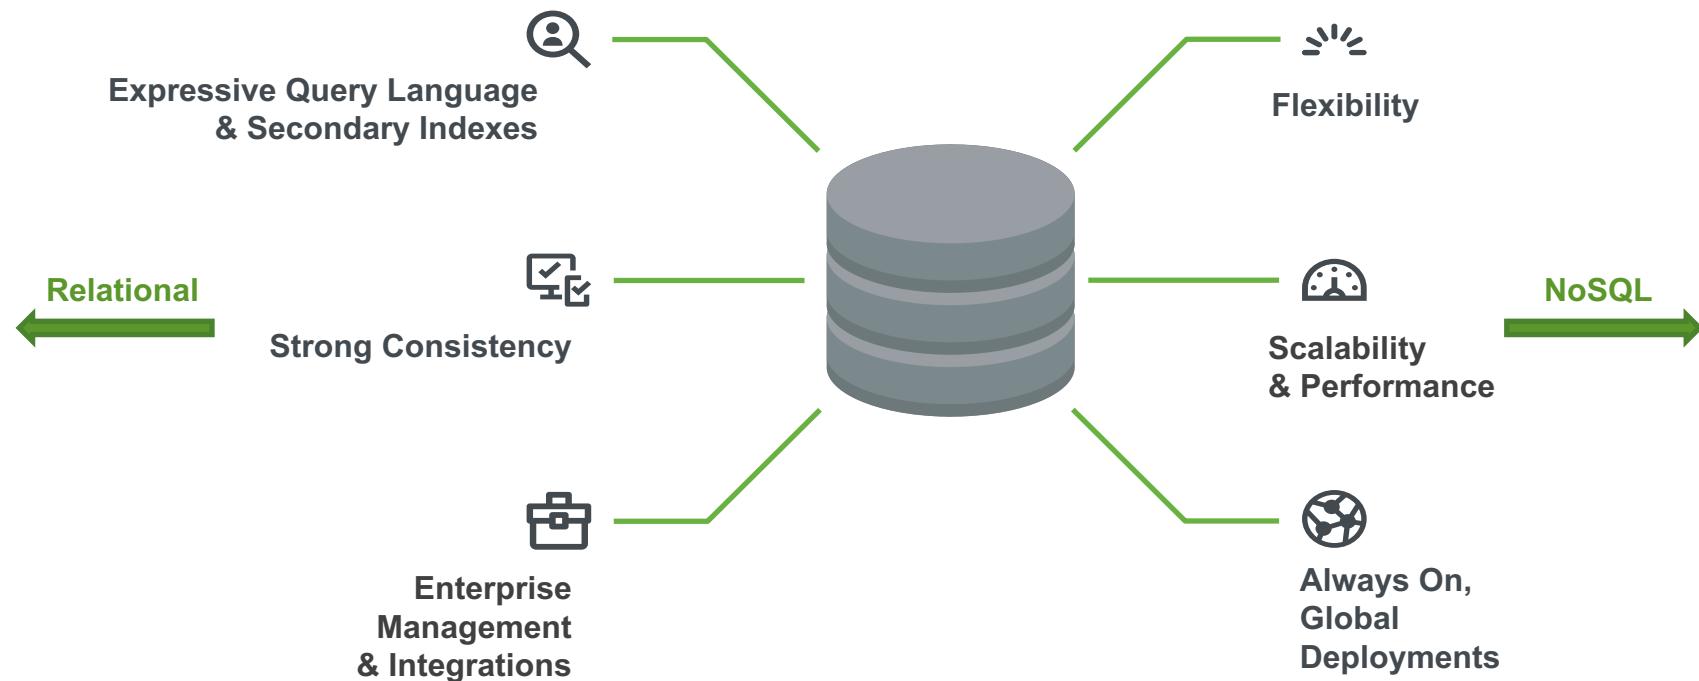


SYBASE®

IBM DB2.



Strengths & Weakness



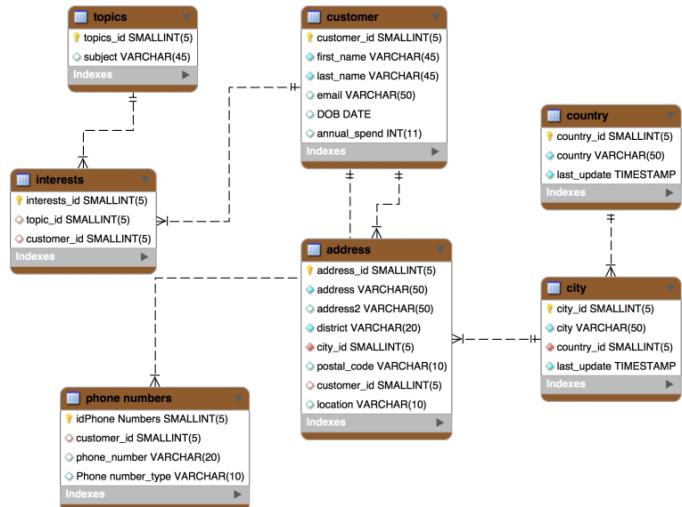


Q1 : Data Model

- All of these data models provide **schema flexibility**. → What do you mean flexibility ?
- The key-value and wide-column data model is opaque in the system - only the primary key can be queried. → is it really need ?
- The document data model has the **broadest applicability**. → Why ?
- The document data model is the **most natural and most productive** because it maps directly to objects in modern object-oriented languages. → more specific for productive ?
- The wide column model provides more granular access to data than the key value model, but less flexibility than the document data model. → why ?



What do you mean ?



Tabular (Relational) Data Model

Related data split across multiple records and tables

```
_id: 12345678
  name: Object
    first: "John"
    second: "Doe"
  address: Array
    0: Object
      location: "work"
      address: Object
        number: 16
        street: "Hatfields"
        city: "London"
        postalCode: "SE1 8DJ"
    1: Object
      location: "home"
      address: Object
    phone: Array
      0: Object
        location: "work"
        number: "+44-1234567890"
      1: Object
      2: Object
        email: "john.doe@mongodb.com"
        dob: 1966-07-30 01:00:00.000
  interests: Array
    0: "Cycling"
    1: "IoT"
```

Document Data Model

Related data contained in a single, rich document



What do you mean ?

row key	columns ...			
	name	email	address	state
jbellis	jonathan	jb@ds.com	123 main	TX
	daria	dh@ds.com	45 2 nd St.	CA
egilmore	name	email		
	eric	eg@ds.com		

Tabular (Relational) Data Model

Related data split across multiple records and tables

```
_id: 12345678
✓ name: Object
  first: "John"
  second: "Doe"
✓ address: Array
  ✓ 0: Object
    location: "work"
    ✓ address: Object
      number: 16
      street: "Hatfields"
      city: "London"
      postalCode: "SE1 8DJ"
  ✓ 1: Object
    location: "home"
    > address: Object
✓ phone: Array
  ✓ 0: Object
    location: "work"
    number: "+44-1234567890"
  > 1: Object
  > 2: Object
    email: "john.doe@mongodb.com"
    dob: 1966-07-30 01:00:00.000
✓ interests: Array
  0: "Cycling"
  1: "IoT"
```

Document Data Model

Related data contained in a single, rich document



Q2 : Query Model

- The biggest difference between non-tabular databases lies in the ability to **query data efficiently**. → ambiguity
- Document databases provide the **richest query functionality**, which allows them to address a wide variety of operational and real-time analytics applications. → compare other data stores
- Key-value stores and wide column stores provide a single means of accessing data: by primary key. This can be fast, but they offer very **limited query functionality** and may impose additional development costs and application-level requirements to support anything more than basic query patterns. → it is not purpose for polygolot ?
- HA & Scales View of Points → utilize the resources ?



What do you mean ?

Expressive Queries

- Find anyone with phone # “1-212...”
- Check if the person with number “555...” is on the “do not call” list

Geospatial

- Find the best offer for the customer at geo coordinates of 42nd St. and 6th Ave

Text Search

- Find all tweets that mention the firm within the last 2 days

Aggregation

- Count and sort number of customers by city, compute min, max, and average spend

Native Binary JSON support

- Add a phone number to Mark Smith’s
- Update just 2 phone numbers out of 10
- Sort on the modified date

JOIN (\$lookup)

- Query for all San Francisco residences, lookup their transactions, and sum the amount by person

Graph queries (\$graphLookup)

- Query for all people within 3 degrees of separation from Mark

{

1212",

1213",

...

MongoDB

```
customer_id : 1,  
first_name : "Mark",  
last_name : "Smith",  
city : "San Francisco",  
phones: [ {  
    number : "1-212-777-  
    type : "work"  
},  
{  
    number : "1-212-777-  
    type : "cell"  
}]
```



What do you mean ?

Fully Indexable

Fully featured secondary indexes

Index Types

- ❖ Primary Index
 - Every Collection has a primary key index
- ❖ Compound Index
 - Index against multiple keys in the document
- ❖ MultiKey Index
 - Index into arrays
- ❖ Text Indexes
 - support for text searches
- ❖ GeoSpatial Indexes
 - 2d & 2dSphere indexes for spatial geometries
- ❖ Hashed Indexes
 - Hashed based values for sharding

Index Features

- ❖ TTL Indexes
 - Single Field indexes, when expired delete the document
- ❖ Unique Indexes
 - Ensures value is not duplicated
- ❖ Partial Indexes
 - Expression based indexes, allowing indexes on subsets of data
- ❖ Case Insensitive Indexes
 - supports text search using case insensitive search
- ❖ Sparse Indexes
 - Only index documents which have the given field



Q3 : Consistency and Transaction Model

- Most applications and development teams expect **strongly consistent** systems. → consistency is depends on the rules ?
- Different consistency models pose **different trade-offs** for applications in the areas of consistency and availability. → what is trade off ?
- MongoDB provides **tunable consistency**, defined at the query level. → what do you mean tunable ?
- Most non-tabular databases provide single record atomicity. This is sufficient for many applications, **but not for all**. MongoDB provides multi-document ACID guarantees, making it easier to address a complete **range of use-cases with a single data platform** → data platform is ambiguity ?



What do you mean ?

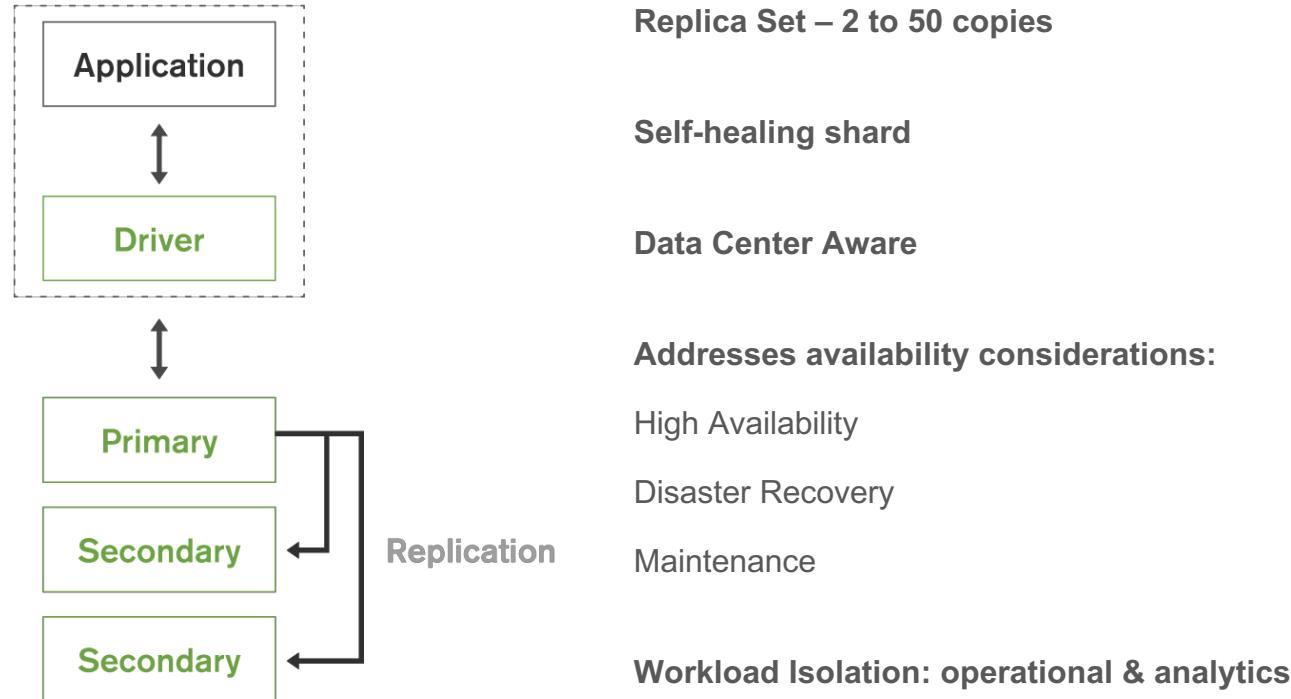
```
with client.start_session() as s:  
    s.start_transaction()  
    try:  
        collection.insert_one(doc1, session=s)  
        collection.insert_one(doc2, session=s)  
        s.commit_transaction()  
    except Exception:  
        s.abort_transaction()
```

Natural for developers

- Idiomatic to the programming language
- Familiar to relational developers
- Simple

*Python. Syntax is subject to change

What do you mean ? High Availability





What do you mean ? Scaling Out

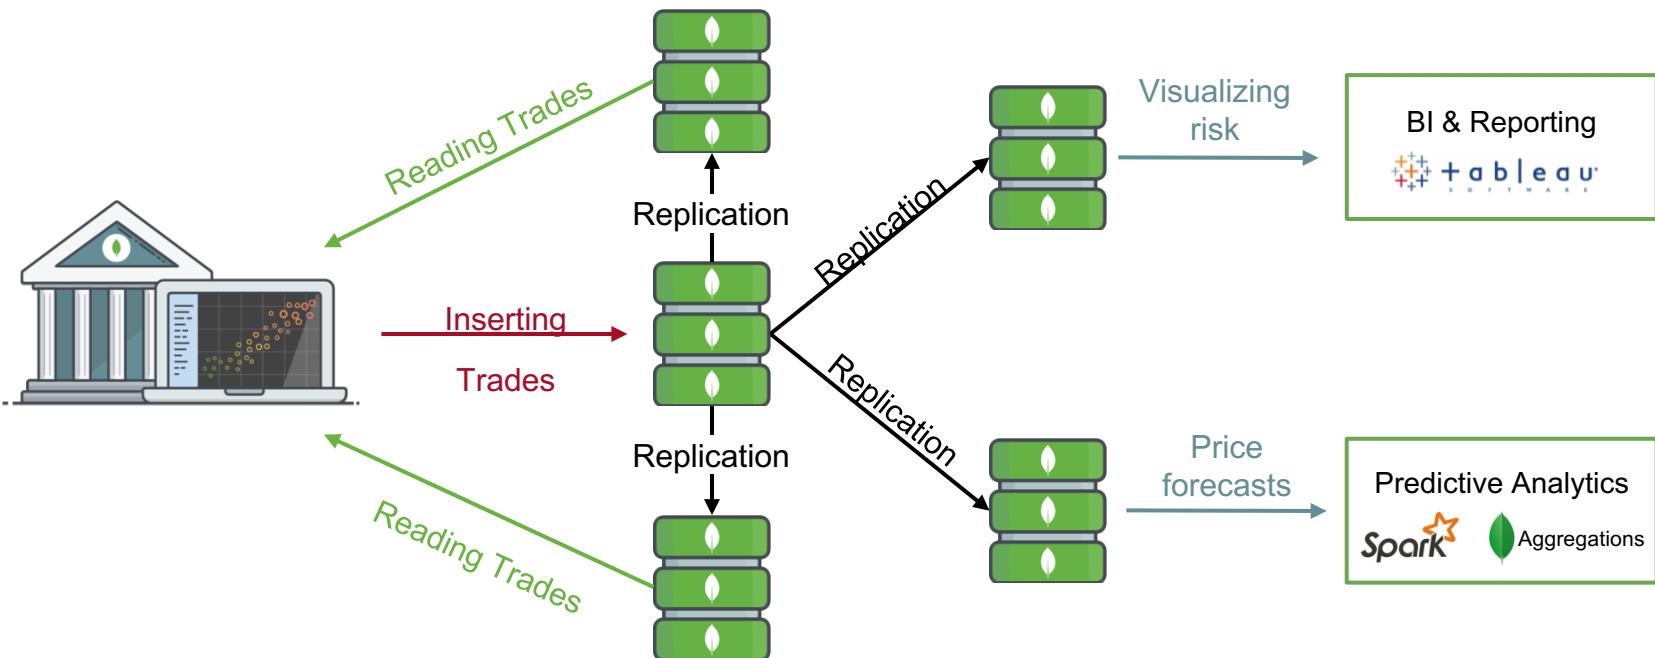


Three types: hash-based, range-based, location-aware

Increase or decrease capacity as you go

Automatic balancing

What do you mean ? Workload Isolation





Q4 : API

- The maturity and functionality of APIs **vary significantly** across non-relational products.
- **MongoDB's idiomatic drivers** minimize onboarding time for new developers and simplify application development.
- **Not all SQL is created equal.** Carefully evaluate the SQL-like APIs offered by non-relational databases to ensure they can meet the needs of your application and developers. → How do you think about that ?
- Is it really important to interact other system / env . ?



Example : SQL JOINS and aggregation

SELECT

```
city,  
SUM(annual_spend) Total_Spend,  
AVG(annual_spend) Average_Spend,  
MAX(annual_spend) Max_Spend,  
COUNT(annual_spend) customers
```

```
FROM (  
    SELECT t1.city, customer.annual_spend  
    FROM customer
```

```
    LEFT JOIN (  
        SELECT address.address_id, city.city,  
               address.customer_id, address.location  
        FROM address LEFT JOIN city  
        ON address.city_id = city.city_id  
    ) AS t1
```

```
    ON  
        (customer.customer_id = t1.customer_id AND  
         t1.location = "home")  
) AS t2
```

```
GROUP BY city;
```

SQL queries have a nested structure

Understanding the outer layers requires understanding the inner ones

So SQL has to be read “inside-out”



Example : Complex queries fast to create, optimize, & maintain

```
db.customers.aggregate([
  {
    $unwind: "$address",
  },
  {
    $match: {"address.location": "home"}
  },
  {
    $group: {
      _id: "$address.city",
      totalSpend: {$sum: "$annualSpend"},
      averageSpend: {$avg: "$annualSpend"},
      maximumSpend: {$max: "$annualSpend"},
      customers: {$sum: 1}
    }
  }
])
```



These “phases” are distinct and easy to understand

They can be thought about in order... no nesting.

MongoDB's aggregation framework has the flexibility you need to get value from your data, but without the complexity and fragility of SQL



Tooling example

aggregations.movies

Documents Aggregations Schema Explain Plan Indexes Validation

Sample pipeline SAVE PIPELINE ... ADD STAGE

Select an operator to construct expressions used in the aggregation pipeline stages. Learn more

\$match Sample of Documents after the \$match stage

```
1 + [{  
2   "_id": ObjectId("573a1398f29313caabcf4fc1"),  
3   "title": "A Turn of the Century Illusionist",  
4   "year": 1899,  
5   "runtime": 1,  
6   "cast": [],  
7   "lastupdated": "2015-09-29 01:21:21.547000000",  
8   "type": "movie",  
9   "directors": []}]
```

```
1 + [{  
2   "_id": ObjectId("573a1395f29313caabce2999"),  
3   "title": "Those Magnificent Men in Their Flying  
        Machines or How I Flew from London...",  
4   "year": 1970,  
5   "runtime": 138,  
6   "released": "1965-06-15T17:00:00.000",  
7   "cast": [],  
8   "poster": "http://ia.media-imdb.com/images/M/M5BNTSMjYzHTY3H1SBH15b15"}]
```

```
1 + [{  
2   "_id": ObjectId("573a1396f29313caabce4fc1"),  
3   "title": "Topgolf Hole",  
4   "year": 1913,  
5   "runtime": 96,  
6   "released": "1913-10-26T17:00:00.000",  
7   "cast": [],  
8   "poster": "http://ia.media-imdb.com/images/M/M5BNTSMjYzHTY3H1SBH15b15"}]
```

```
1 + [{  
2   "_id": ObjectId("573a13bfdf29313caabcf41f0"),  
3   "title": "Ella Lola, la Trilby",  
4   "year": 1899,  
5   "runtime": 96,  
6   "released": "1900-01-01T00:00:00.000",  
7   "cast": [],  
8   "plot": "Dancer Ella Lola dances a routine based on the famous character of \"Trilby\".",  
9   "fullplot": "Dancer Ella Lola dances a routine based on the famous character of \"Trilby\"."}]
```

\$project Sample of Documents after the \$project stage

```
1 + [{  
2   "title": "A Turn of the Century Illusionist",  
3   "year": 1899,  
4   "directors": []}]
```

```
1 + [{  
2   "title": "Those Magnificent Men in Their Flying  
        Machines or How I Flew from London...",  
3   "year": 1970,  
4   "cast": []}]
```

```
1 + [{  
2   "title": "Topgolf Hole",  
3   "year": 1913,  
4   "cast": []}]
```

```
1 + [{  
2   "title": "Ella Lola, la Trilby",  
3   "year": 1899,  
4   "cast": []}]
```

\$sort Sample of Documents after the \$sort stage

```
1 + [{  
2   "year": 2012, "title": "Life of Pi"},  
3   {"year": 2010, "title": "Babies"}]
```

```
1 + [{  
2   "year": 2010, "title": "Babies"},  
3   {"year": 2012, "title": "Life of Pi"}]
```

```
1 + [{  
2   "year": 2010, "title": "The Secret World of Arrietty"},  
3   {"year": 2010, "title": "Babies"}]
```

- Easily build aggregation pipelines for data processing and analytics
 - Add new pipeline stages with code skeletons and auto-completion
 - See a preview of documents as they are transformed in each stage
 - Drag and drop to re-order pipeline stages
 - Save pipelines in Compass or export to native code

Introduced in MongoDB Compass 1.14



Tooling Example : Export to Language



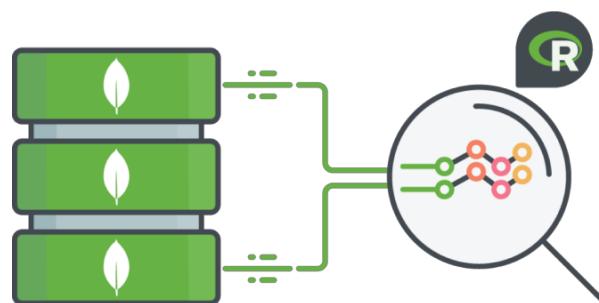
- Simplify query creation: build a query in MongoDB Compass, then click to export it to native code
- Supports complex, multi-parameter queries and aggregation pipelines
- Initially supported: C#/.NET, Java, JavaScript, Python, and mongo shell syntax
 - More languages to come in future releases

*Introduced in MongoDB Compass
1.15*



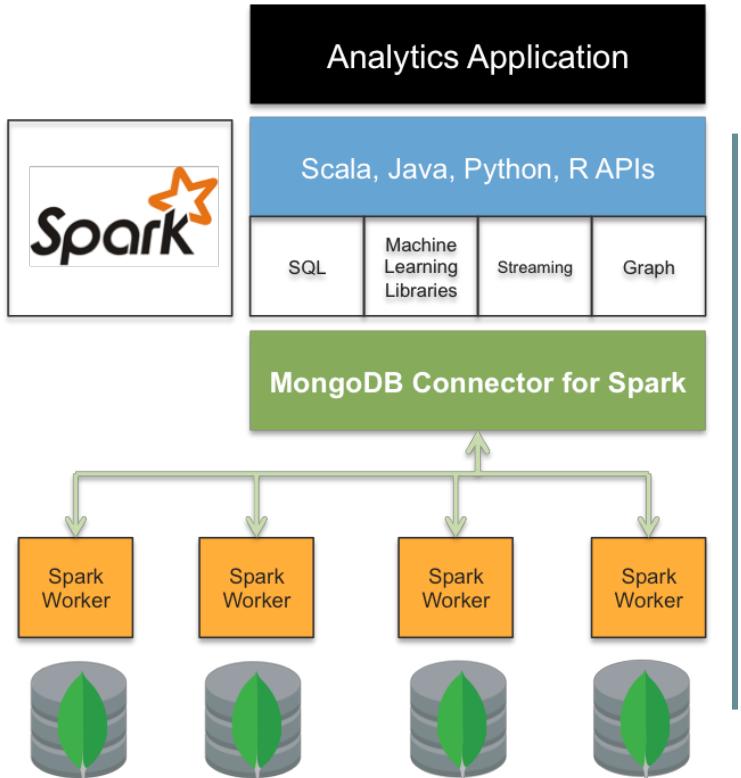
R Driver for MongoDB

Recommended MongoDB R driver for data scientists, developers & statisticians



- **Idiomatic**, native language access to the database
- **MongoDB read & write concerns** to control data consistency & durability
- **Data security** with enterprise authentication mechanisms
- **Advanced BSON data types**, e.g., Decimal 128 for high precision scientific & financial analysis

MongoDB Connector for Apache Spark



- Native Scala connector, certified by Databricks
- Exposes all Spark APIs & libraries
- Efficient data filtering with predicate pushdown, secondary indexes, & in-database aggregations
- Locality awareness to reduce data movement

"We reduced 100+ lines of integration code to just a single line after moving to the MongoDB Spark connector."

- Early Access Tester, Multi-National Banking Group



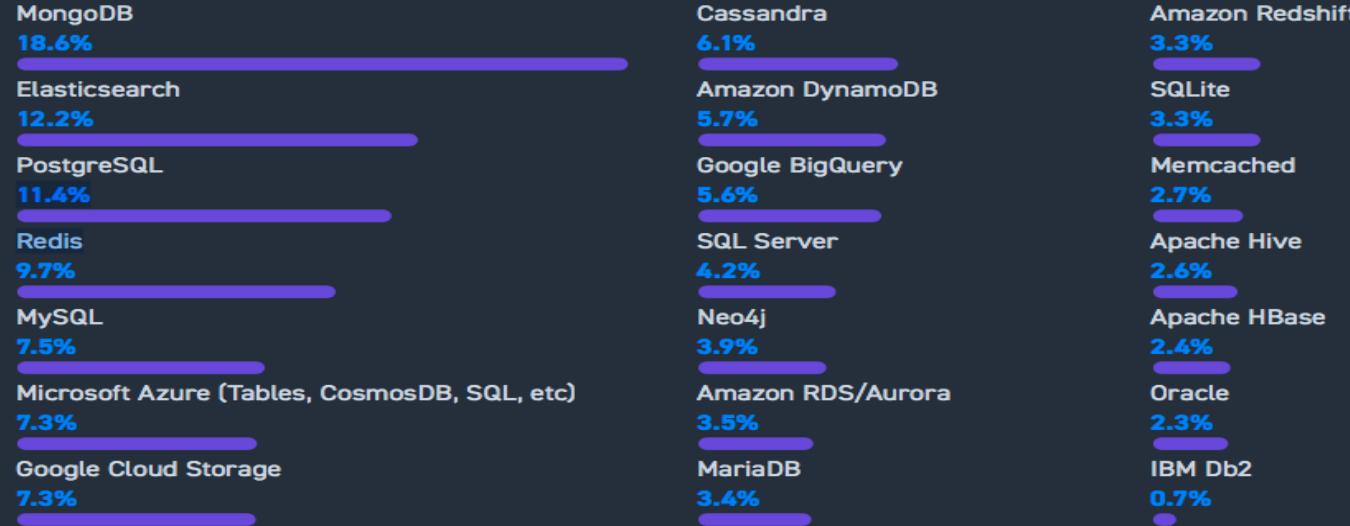
Q5 : Commercial Support, Community Strength, Freedom from Lock-In

- **Community size and commercial strength** is an important part of evaluating non-relational databases.
- MongoDB is one of the very few non-relational database providers to be a publicly traded company; it has the largest and most active community; support teams spread across the world providing 24x7 coverage; user-groups in most major cities; and extensive documentation.
- MongoDB is **available** to run on your own infrastructure, or as a fully managed cloud service on all of the leading public cloud platforms.

What do you mean ? Stack Overflow Developer Survey 2018 , 2019



Most Wanted Databases



The Forrester Wave™: Big Data NoSQL, Q1 2019

THE FORRESTER WAVE™

Big Data NoSQL

Q1 2019



The Largest Ecosystem

45,000,000+

MongoDB Downloads

850,000+

Online Education Students

100,000+

MongoDB Atlas Clusters

55,000+

MongoDB User Group Members

1,000+

Technology and Services Partners

5,700+

Customers Across All Industries

DB-Engines.com

346 systems in ranking, October 2018

Rank	Oct 2018	Sep 2018	Oct 2017	DBMS	Database Model	Score		
						Oct 2018	Sep 2018	Oct 2017
1.	1.	1.	1.	Oracle 	Relational DBMS	1319.27	+10.15	-29.54
2.	2.	2.	2.	MySQL 	Relational DBMS	1178.12	-2.36	-120.71
3.	3.	3.	3.	Microsoft SQL Server 	Relational DBMS	1058.33	+7.05	-151.99
4.	4.	4.	4.	PostgreSQL 	Relational DBMS	419.39	+12.97	+46.12
5.	5.	5.	5.	MongoDB 	Document store	363.19	+4.39	+33.79
6.	6.	6.	6.	DB2 	Relational DBMS	179.69	-1.38	-14.90
7.	↑ 8.	↑ 9.	9.	Redis 	Key-value store	145.29	+4.35	+23.24
8.	↓ 7.	↑ 10.	10.	Elasticsearch 	Search engine	142.33	-0.28	+22.09
9.	9.	↓ 7.	7.	Microsoft Access	Relational DBMS	136.80	+3.41	+7.35
10.	10.	↓ 8.	8.	Cassandra 	Wide column store	123.39	+3.83	-1.40

Method of calculating the scores of the DB-Engines Ranking

The DB-Engines Ranking is a list of database management systems ranked by their current popularity. We measure the popularity of a system by using the following parameters:

- **Number of mentions of the system on websites**, measured as number of results in search engines queries. At the moment, we use [Google](#), [Bing](#) and [Yandex](#) for this measurement. In order to count only relevant results, we are searching for <system name> together with the term database, e.g. "Oracle" and "database".
- **General interest in the system**. For this measurement, we use the frequency of searches in [Google Trends](#).
- **Frequency of technical discussions about the system**. We use the number of related questions and the number of interested users on the well-known IT-related Q&A sites [Stack Overflow](#) and [DBA Stack Exchange](#).
- **Number of job offers, in which the system is mentioned**. We use the number of offers on the leading job search engines [Indeed](#) and [Simply Hired](#).
- **Number of profiles in professional networks, in which the system is mentioned**. We use the internationally most popular professional networks [LinkedIn](#) and [Upwork](#).
- **Relevance in social networks**. We count the number of [Twitter](#) tweets, in which the system is mentioned.



Why bet on us?



Ease and flexibility of the document model led to massive early adoption. We are the **world's most popular modern database**.



First database company to go public in over 20 years. We've raised **over a half a billion dollars** to invest in our business.



Uniquely positioned as the **only company who can credibly offer a modern “general purpose” database**.



Appendix : MongoDB Atlas & Proof of Points





MongoDB in Database World

Intelligent Operational Data Platform



**Best way to
work with data**



**Intelligently put data
where you need it**



**Freedom
to run anywhere**

MongoDB Atlas — 글로벌 클라우드 매니즈드 데이터베이스 서비스

Self-service & elastic <ul style="list-style-type: none">- 필요시 배포, 수정 및 업그레이드는 베스트 Practice 기반으로 운영 자동화 제공- 자동화 된 데이터베이스 유지 관리 몇 번의 클릭 또는 API 호출로 확대, 축소 또는 축소 구성의 자동화	Global & cloud-agnostic <ul style="list-style-type: none">- AWS 여러 지역에서 사용 가능- 어디에서나 읽기 / 쓰기가 가능한 글로벌 클러스터를 통한 배포 및 다중 리전을 통한 고가용성 제공- 클라우드 환경 기반에서 손쉽게 마이그레이션 가능	Enterprise-grade security & SLAs <ul style="list-style-type: none">- 네트워크 격리, VPC 피어링, 종단 간 암호화 및 역할 기반 액세스 제어- 암호화 키 관리, LDAP 통합, 상세 데이터베이스 감사- SOC 2 / 개인 정보 보호/ HIPAA SLA로 안정성 보장
Comprehensive monitoring <ul style="list-style-type: none">- 전 100 개 이상의 메트릭을 통한 KPI에 대한 alerting- 실시간 성능 조언 및 성능 관리자- 모니터링 대시 보드와 통합 할 수 있는 API	Managed backup <ul style="list-style-type: none">- 특정 시점의 데이터 복구- 쿼리 가능한 백업 스냅 샷- 샤크환경내 일관된 스냅 샷- 클라우드 데이터 이동성	Stitch: Serverless platform services <ul style="list-style-type: none">- 백엔드 로직, 서비스 통합 및 API를 위한 단순한 서비스 기능- 간단한 필드 레벨 액세스 규칙으로 보호되는 프론트 엔드에서의 데이터베이스 액세스- 변경 사항에 실시간으로 반응하는 데이터베이스 및 인증 트리거

MongoDB Atlas을 사용할 수 있는 AWS Region

MongoDB Atlas is currently available in 16 AWS regions.

Create a **free tier cluster** by selecting a region with **FREE TIER AVAILABLE** and choosing the **M0** cluster tier below.

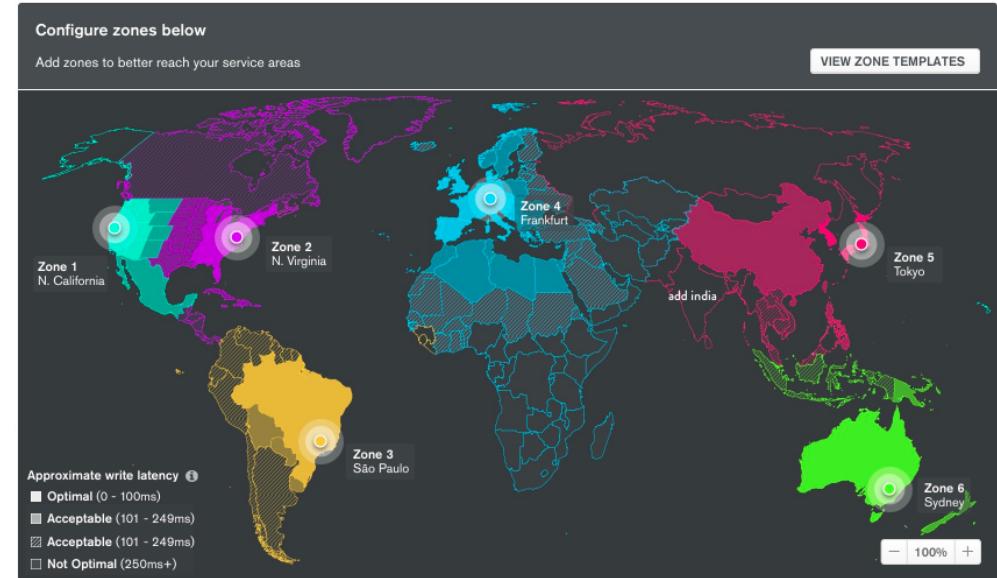
★ recommended region i

NORTH AMERICA	EUROPE	AUSTRALIA	ASIA
N. Virginia (us-east-1) ★ FREE TIER AVAILABLE	Stockholm (eu-north-1) ★	Sydney (ap-southeast-2) ★	Tokyo (ap-northeast-1) ★
Ohio (us-east-2) ★	Ireland (eu-west-1) ★ FREE TIER AVAILABLE		Seoul (ap-northeast-2)
N. California (us-west-1)	London (eu-west-2) ★		Singapore (ap-southeast-1) ★ FREE TIER AVAILABLE
Oregon (us-west-2) ★ FREE TIER AVAILABLE	Paris (eu-west-3) ★		Mumbai (ap-south-1) FREE TIER AVAILABLE
Montreal (ca-central-1)	Frankfurt (eu-central-1) ★ FREE TIER AVAILABLE		
SOUTH AMERICA	Sao Paulo (sa-east-1)		

글로벌 클러스터

하나 이상의 클라우드 영역으로 구성되고, 지리적으로 분산 된 여러 영역에 걸쳐 완전 자동화 된 데이터베이스 배포

- 분산 응용 프로그램에 대해 한 자리 수 (밀리 초) 대기 시간을 제공하기 위해 로컬로 읽고 쓰기
- 사전 제작 된 영역 템플릿을 사용하여 손쉽게 배포하거나 사용하기 쉬운 시각적 인터페이스에서 클라우드 영역을 선택하여 자체 영역을 구축하십시오.



다중 Region을 통한 replication set 구성

- 드물게 발생할 수 있는 전체 지역 장애에 대해
다중 리전 구성으로 고 가용성을 보장하십시오.
- Voting에 참여하지 않는 읽기 전용 노드 구성을
통해 (지리적으로 배포) 로컬 읽기 성능 향상할
수 있습니다.
- AWS Region 기반에서 구성, 배포 가능 합니다.



다중 Region을 통한 replication set 구성

- Electable Node에 대한 multi region 구성으로 single region에 대한 outage를 대비 할 수 있다.
- Read Only Node를 통해 (Local Node) 구성을 통해 latency 문제를 해결 할 수 있다.
- Analytic Node를 통해 (Local Node) 구성을 통해 latency 문제를 해결 할 수 있다.

Electable nodes for high availability
Configure 3, 5, or 7 nodes across multiple regions to better withstand data center outages

Region	Nodes
 Seoul (ap-northeast-2)	  3

+ Add a region

Read-only nodes for optimal local reads
Add replicas in additional regions to optimize for local reads in any of your service areas

Region	Nodes
 Seoul (ap-northeast-2)	 1 
 Tokyo (ap-northeast-1)	 1 

+ Add a region

Analytics nodes for workload isolation
Isolate queries on read-only nodes that will not contend with your operational workload

Region	Nodes
 Seoul (ap-northeast-2)	 1 

실시간 성능 모니터링

실시간 성능 패널은 클러스터에서 발생한 상황을 실시간으로 확인하고 문제를 신속하게 진단을 통해 탁월한 통찰력을 데이터베이스에 제공합니다.

디스플레이 항목

- Operations
- Read/Writes
- Network In/Out
- Memory
- Hottest Collections
- Slowest Operations



Fine-grained 모니터링 & 경보

- 모니터링 및 경고(alerting)는 클러스터 데이터베이스 및 서버 사용 상태에 대한 전체 메트릭을 제공합니다.
- 데이터베이스 작업 또는 서버 사용량이 클러스터 성능에 영향을 미치는 정의된 임계 값에 도달하면 자동 알림
- 자동화 된 경고 기능과 MongoDB Atlas의 유연한 확장 및 축소 옵션을 결합하여 데이터베이스 지원 응용 프로그램을 항상 수행 할 수 있습니다.



Fine-grained 모니터링 & 경보

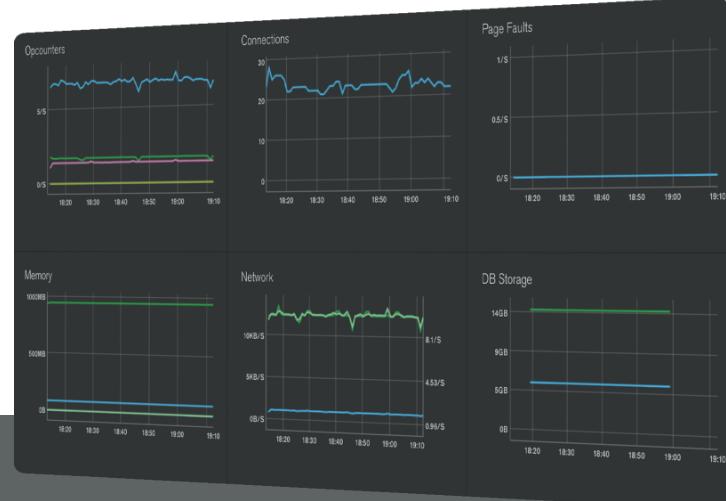
TOGGLE CHARTS

MongoDB Metrics

- Asserts
- Cache Activity
- Cache Usage
- Connections
- Cursors
- + DB Storage
- Document Metrics
- + Memory
- Network
- Opcounters
- + Opcounters - Repl
- Operation Execution Times
- + Oplog GB/Hour
- Page Faults
- + Query Executor
- Query Targeting
- + Queues
- + Replication Headroom
- + Replication Lag
- + Replication Oplog Window
- Scan And Order
- Tickets Available

Hardware Metrics

- + Disk IOPS
- + Disk Latency
- + Disk Space Free
- + Disk Space Percent Free
- + Disk Space Used
- + Normalized Process CPU
- + Normalized System CPU
- + Process CPU
- + System CPU
- Util %



성능 어드바이저

MongoDB Atlas 클러스터에서 사용할 수 있는 Always-on Performance Advisor는 성능 향상에 도움이 되는 자동화 된 권장 사항을 제공합니다.

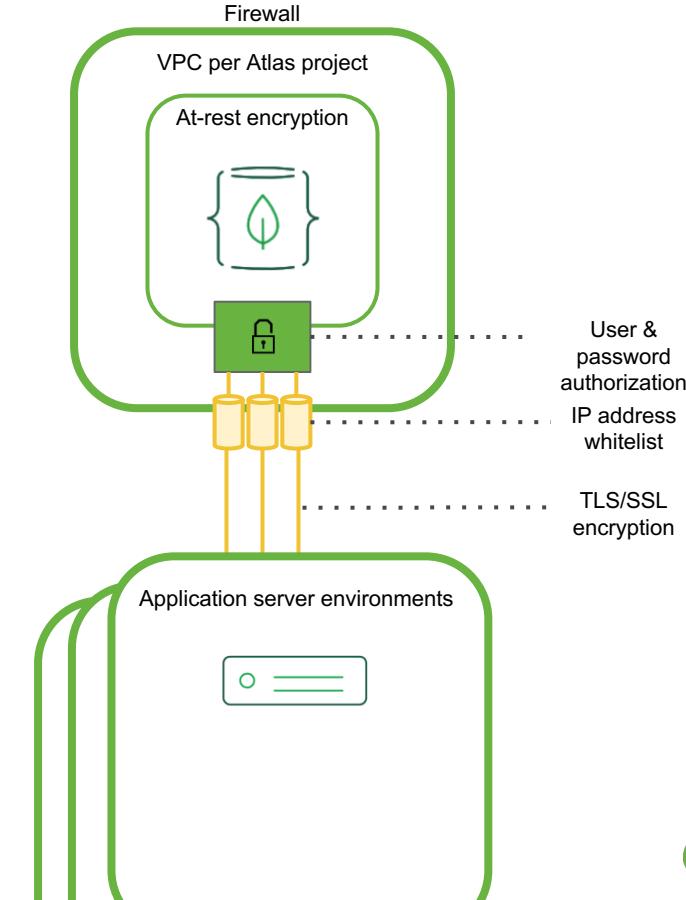
- 느리게 실행되는 쿼리에 대한 관련 통계를 보고 자동화 된 인덱스 제안 수신
- 클러스터의 기존 인덱스 보기
- The Performance Advisor는 성능 오버 헤드 없이 언제든지 실행할 수 있습니다.

The screenshot shows the MongoDB Atlas Performance Advisor interface. On the left, there's a sidebar with project navigation (Clusters, Stitch Apps, Alerts, Backup, Users, Settings, Orgs and Projects, Docs, Support) and a context dropdown set to 'andrewAtlas'. The main area has tabs for 'EXISTING INDEXES' and 'RECOMMENDED INDEXES'. Under 'RECOMMENDED INDEXES', a card for 'CREATE INDEX callLetters: 1' is shown with a 'More Info' button. Below it, under 'SAMPLE QUERIES IMPROVED BY THIS INDEX', there's a 'FIND callLetters: "UDZH"' section with a 'SHOW MORE' button. At the bottom, there's a 'METRICS FOR THESE QUERIES' section with a table:

QUERY INEFFICIENCY SCORE	EXECUTION COUNT	AVERAGE EXECUTION TIME
12045768	101	447554 MS

보안 : 고객 전용 VPC

- 네트워크는 기본적으로 외부와 연결되어 있지 않습니다.
- IP 주소는 인바운드 네트워크 트래픽을 위해 명시적으로 허용 목록에 있어야합니다.
- 데이터 암호화 (in- flight) - 모든 네트워크 트래픽에 사용되는 TLS / SSL
- 데이터베이스에 연결하는 데 필요한 사용자 이름과 암호. 구성 가능한 권한
- 암호화 된 저장소(at rest) 볼륨에서의 데이터 암호화



보안 : VPC Peering

- **VPC Peering** — Atlas 클러스터가 포함 된 VPC를
사설 네트워크를 통해 연결하는 응용 프로그램 서버가
포함 된 VPC에 피어 투 피하십시오.

Peering Connection AWS M10 & ABOVE ONLY

Your Application VPC

Account ID
You can find this in "My Account" in your AWS Console. [Learn More](#)

VPC ID
You can find this in your list of VPCs in the VPC dashboard in your AWS account. [Learn More](#)

VPC CIDR
This CIDR block cannot overlap with your Atlas VPC CIDR block shown below or the peering will fail. [Learn More](#)

Application VPC Region
supports VPC peering with other AWS VPCs in both the same and different regions.

Your Atlas VPC

VPC CIDR
A custom Atlas CIDR block must be at least a /24 and at most a /27. You cannot modify the CIDR block if you have an existing cluster.

Atlas VPC Region
Uncheck the checkbox to select an Atlas VPC region different from that of your application's VPC.

Cancel **Initiate Peering**

